

Diego Avina-Escobedo

Shirley Moore

CS 4375

HW 5 Report:

Task 1: Implementing mmap() and munmap() with Lazy Allocation:

a) Results of Running Private Program:

I successfully implemented mmap() and munmap() following the instructions in the “Anonymous mmap() and munmap()” handout. The private program uses mmap() to allocate a region of virtual memory and munmap() to deallocate it. However, before modifying usertrap() in kernel/trap.c for part b, the program aborted due to a fault that was not handled

```
..          1 1 1024
README     2 2 2226
cat        2 3 23856
echo       2 4 22672
forktest   2 5 13416
grep       2 6 26992
init       2 7 23496
kill       2 8 22616
ln         2 9 22472
ls         2 10 26024
mkdir      2 11 22736
rm         2 12 22728
sh         2 13 40736
stressfs   2 14 23712
usertests  2 15 150432
grind      2 16 37208
wc         2 17 24816
zombie     2 18 21984
private    2 19 23840
console    3 20 0
$ private
usertrap(): unexpected scause 0x0000000000000002 pid=4
             sepc=0x0000000000000028 stval=0x0000000000000000
```

b) Fixing usertrap() for Private Program:

To address the fault issue, I modified `usertrap()` in `kernel/trap.c`. Specifically, I added checks for `scause` being either a load fault (13) or a store fault (15) and ensured that the fault address falls within a mapped memory region with the correct permissions. Additionally, I allocated a physical memory frame using `kalloc()`, mapped it into the process's page table using `mappages()`, and rounded the fault address down to a page boundary. After these modifications, the private program ran successfully.

```
void usertrap(void) {
    int which_dev = 0;
    struct proc *p = myproc();

    if ((r_sstatus() & SSTATUS_SPP) != 0)
        panic("usertrap: not from user mode");

    w_stvec((uint64)kernelvec);
    p->trapframe->epc = r_sepc();

    if (r_scause() == 8) {
        // Handle system call
        if (p->killed) exit(-1);
        p->trapframe->epc += 4;
        intr_on();
        syscall();
    } else if ((which_dev = devintr()) != 0) {
        // Handle device interrupts
    } else if (r_scause() == 13 || r_scause() == 15) {
        // Handle page fault
        if (!is_valid_address(r_stval(), p)) {
            p->killed = 1;
            exit(-1);
        }
        void *physical_mem = kalloc();
        if (!physical_mem) {
            printf("usertrap(): out of memory\n");
            p->killed = 1;
            exit(-1);
        }
        if (mappages(p->pagetable, PGROUNDDOWN(r_stval()), PGSIZE, (uint64)physical_mem, PTE_R | PTE_W | PTE_X | PTE_U) < 0) {
            kfree(physical_mem);
            printf("usertrap(): mappages failed\n");
            p->killed = 1;
            exit(-1);
        }
    } else {
        // Handle unexpected traps
        printf("usertrap(): unexpected scause %p pid=%d\n", r_scause(), p->pid);
    }
}
```

77,12

20%

c) Handling `munmap()` in `freeproc()`

Initially, commenting out the call to `munmap` in `private.c` and running the program caused a kernel panic. This was because the physical memory for the mapped memory region was not freed in `freeproc()`. I added the provided code in the handout to `freeproc()`, preventing the kernel panic.

```
hart 2 starting
hart 1 starting
init: starting sh
$ private
exec private failed
```

```
xv6 kernel is booting
```

```
hart 2 starting
hart 1 starting
init: starting sh
$ private
total = 55
$ █
```

Task 2: Modifying fork() for Shared Memory Inheritance

a) Difference between uvmcopy() and uvmcopyshared()

For this part of task 2, I made the necessary changes to `uvmcopy()` in `vm.c` and added `uvmcopyshared()`. The difference between the two lies on how they handle private and shared mappings. `Uvmcopy()` copies the memory region as is, while `uvmcopyshared()` ensures that changes made by any process in a shared mapping are reflected across all processes sharing that mapping.

b) Modifying fork() for inheriting Shared Memory

I added code to `fork()` in `proc.c` to copy the memory region table from parent to child, ensuring proper inheritance of both private and shared mappings. The code works for both `PRIVATE` and `SHARED` regions, preserving the isolation of private mappings and allowing shared mappings to reflect changes across processes.

c) Testing with `prodcons1.c` and `prodcons2.c`

I compiled the modified code and tested with prodcons1.c and prodcons2.c. The results were as expected, with prodcons1.c showing a total of 55 and prodcons2.c showing a total of 0.

```
xv6 kernel is booting

hart 2 starting
hart 1 starting
init: starting sh
$ prodcons1
total = 55
$ prodcons2
total = 0
$ █
```

Task 3: Handling Multiple Pages in prodcons3.c

a) Incorrect Results in prodcons3.c

Running prodcons4.c, which maps three pages, initially produced incorrect results.

The trap handler in usertrap() needed modifications to properly handle situations where a process other than the one that originally mapped a shared memory region is the first to write to it.

b) Extra Credit: Handling Shared Memory Writes

I modified the trap handler in usertrap() to handle situations where a different process writes to a shared memory region first. The code now correctly inserts the new mapping for the allocated physical page into the page table for all processes in the family that have the shared memory region mapped.

```
$ prodcons3
total = 673720320
$ █
```

Summary:

Through this assignment, I gained a deeper understanding of Linux memory mapping concepts, page tables, and shared memory regions. Implementing `mmap()` and `munmap()`, along with modifying `for()` for shared memory inheritance, provided valuable insights into xv6's memory management.