

Diego Jared Avina

Dr. Moore

CS4375-13948 Fall 2023

### HW 3: Priority-based Scheduler for xv6

#### Task 1: Modify the ps Command:

In task 1, we were told to modified the **ps** command to print the priority of each process. This mainly involved changing the **ps.c** file located inside the **user** directory. In order to retrieve the priority for each process, we created two new methods **getpriority** and **setpriority**.

#### Results of ps command task 1

```
init: starting sh
$ pexec 10 ps
pid      state  size   ppid   name    priority
1        sleeping 12288   0      init     0
2        sleeping 16384   1      sh       0
3        sleeping 12288   2      pexec    10
4        running  12288   3      ps       0
```

#### Task 2: Add a readytime field and Age Calculation

In task 2, we needed to add a new field named **readytime** to the **struct proc** and modified the **ps** command located inside the **user** directory to print a process's age. In order to calculate a process's age we updated the **readytime** field whenever a process became **RUNNABLE**.

The age calculation is done as follows:

**Age = current time (uptime()) – process.readytime;**

**Results:**

```

init: starting sh
pe$ xec 10 ps
pid      state   size    ppid    name      priority    cputime age
1        sleeping 12288    0       init       0           0
2        sleeping 16384    1       sh          0           0
3        sleeping 12288    2       pexec      10          0
4        running  12288    3       ps          0           0      -84214978
$ OEMU: Terminated

```

In this case, for my results, I only get one have and I get a negative number. Im not sure what failed in my code because every time I ran the readytime by itself it returned that same number but positive.

### Task 3: Implement Priority-Based Scheduler

In task 3 we implement a priority-based scheduler. The primary changes were made inside the **proc.c** file, more specifically the **scheduler** function; there were a few changes in other files like **param.h**. After running a few tests in the programs using the priority scheduler, we observed that processes with higher priorities indeed run first.

#### What I learned:

Implementing a priority-based scheduler helps in prioritizing tasks based on their importance. This task has shown us the significance of process priorities in managing system resources efficiently.

#### Difficulties:

I encountered multiple challenges going making the scheduler “work”. I had to carefully debug and understand the xv6 codebase.

### Task 4: Adding Aging to the Priority-Based Scheduler

In Task 4, we added aging to our priority-based scheduler. Our aging policy includes periodically increasing the priority of all processes in the system to prevent lower-priority processes from being starved.

### Difficulties:

Implementing the aging policy required a lot of dedication and coordination with the current scheduler code I had. Debugging and thorough testing were necessary to ensure that some of the code worked correctly.

```
xv6 kernel is booting

init: starting sh
$ pexec 5 matmul & matmul 10 &
$ Time: 14 ticks
Time: 66 ticks
pexec 10 ps
pid      state  size  ppid  name    priority  cputime age
1        sleeping 12288  0     init    0         0
2        sleeping 16384  1     sh      0         0
8        sleeping 12288  2     pexec   10        0
9        running  12288  8     ps      0         0      33751
$
```